

Learning with misspecified models: Overconfidence and Stereotypes

Jimena Galindo

October 08, 2023

Abstract

TBW

1 Introduction

1.1 Related literature

2 Theoretical Framework

I consider multiple theories of belief updating that have been proposed in the literature and that are able to rationalize the prevalence of misspecified beliefs. I focus on two different frameworks and develop a unifying example that allows me to compare the predictions of all the theories. In this section, I outline the theories within each of the frameworks.

2.1 Framework 1

An agent is of type $\theta \in \Theta$ and faces an unknown exogenous state ω drawn from some density f over Ω . The agent knows the distribution of ω but not its realized value. His prior

belief about his type is $p_0(\theta)$ and his belief about the state, $p_0(\omega)$, coincides with the true distribution. Let the agent's true type be θ^* , and the realized state be ω^* .

An agent has a *misspecified* belief if the prior assigns probability zero to their true type. Furthermore, the agent is *dogmatic* if he holds a degenerate belief that places probability one on being a particular type, $\hat{\theta}$. An agent can be dogmatic and misspecified. If the agent is both, $\hat{\theta} \neq \theta^*$ and $p_0(\hat{\theta}) = 1$.

The agent chooses an action $a \in A$ and observes a noisy outcome $h \in H$. The outcome is a function of the agent's type, the state, and the action. In particular $h = h(\theta^*, \omega^*, a) + \varepsilon$ with $h(\cdot)$ increasing in both θ^* and ω^* , and such that conditional on a pair of parameters (θ, ω) , there is a unique optimal action. $\varepsilon \sim N(0, \sigma)$ is noise in the output.

After observing the outcome, the agent updates his beliefs about θ and ω using some algorithm and moves on to the next period. He repeats this process infinitely many times. I make the simplifying assumption that the agent is myopic and chooses the action that maximizes the payoff in each period. This assumption simplifies the analysis and plays a role in whether an agent who updates their beliefs using Bayes rule would learn the truth or not. The behavior of a forward-looking agent is discussed in the conclusion.

A key notion in this setting is that of a self-defeating equilibrium¹. A *self-defeating equilibrium* is a belief and action pair such that the agent's belief about their type is misspecified and the outcome generated by the action is consistent with the misspecified belief. This means that the average outcome under the true type and the true state equals the average output the agent expects under the misspecified belief. The agent's belief is said to be *stable* when this happens.

Within this framework, I consider two nested theories of belief updating: the first one is a dogmatic modeler from Heidhues et al. [2018], The second one is a switcher, as in Ba [2023]. The dogmatic modeler can be seen as a switcher with an infinitely sticky initial belief and this is the sense in which the two are nested. Both theories produce different predictions

¹This notion is an adaptation of the Berk-Nash equilibrium in @Esponda2016 to this setting with only one agent

about the agent's behavior. I discuss both in what follows.

2.1.1 The Dogmatic Modeler

A dogmatic agent does not update their beliefs about θ , instead, he holds a degenerate belief that places probability on $\hat{\theta}$, which is potentially misspecified. In this case, no matter how much information he gathers against being of type $\hat{\theta}$, he will not update his beliefs about θ . Any discrepancies between the observed outcomes and his believed type are incorporated using the Bayes rule to update their beliefs about ω . Heidhues et al. [2018] show that, under certain assumptions on the per-period utility,² a dogmatic modeler will inevitably fall into a self-defeating equilibrium. The equilibrium will be such that the outcomes they observe reinforce their belief on ω in such a way that as $t \rightarrow \infty$ the agent will be sure that the state is some ω' consistent with their believed type and the observed data. In other words, they will be in a self-defeating equilibrium with a stable belief that places probability one on the incorrect parameters $(\hat{\theta}, \omega_\infty)$

The mechanism by which the dogmatic agent falls into the self-defeating equilibrium is the following: Suppose the agent holds the misspecified belief that they are type $\hat{\theta} > \theta^*$. For any prior over ω , the agent will be disappointed by the outcome. He expected a gain of $h(\hat{\theta}, \mathbb{E}(\omega), a)$ but instead observes $h(\hat{\theta}, \omega^*, a)$. There are two possible sources for the disappointment, the first is that the realized state is lower than the expected state. The second source is that the agent is of type θ^* and therefore, for all possible states, his gain will be lower than what he expected. Because the agent is dogmatic, he will not update his beliefs about θ and as a consequence will attribute the disappointment to the state being lower than expected. He will continue to update in this way until he converges to a belief about ω that is stable. Such a belief will explain the observed utility perfectly and allow the agent to rationalize his dogmatic belief about θ . Under the assumptions of Heidhues et al.

²The assumptions are that u is twice continuously differentiable with: (i) $u_{ee} < 0$ and $u_e(e\theta, \omega) > 0 > u_e(\bar{e}, \theta, \omega)$, (ii) $u_\theta, u_\omega > 0$ and (iii) $u_{e\theta} < 0$ and $u_{e\omega} > 0$. The direction of the derivatives is a normalization and the results would hold even when the signs are reversed.

[2018], there is a unique value of ω at which the belief is stable, I will refer to such value as ω_∞ . This mechanism is further illustrated in Example 1.

Example 1: Set $A = \Omega$ and $H = [0, \infty)$ and consider a student with intrinsic ability $\theta^* \geq 0$ who faces a grading procedure ω^* that is unknown to them. However, they know that a higher ω^* is more likely to yield a higher grade. In particular, assume the grade is given by $(\theta^* + a)\omega^*$.

The student must choose an effort level a , which determines their grade. For whatever the chosen effort level, the agent must pay a cost $c(a) = \frac{1}{2}a^2$. And he repeats this process for infinitely many periods. Assume also that the student's prior is such that $\mathbb{E}[\omega] = \omega^*$ and he is dogmatic about being of type $\hat{\theta} > \theta^*$.³ Therefore, the student's payoff in period t is given by

$$u_t(a_t; \theta^*, \omega^*) = (\theta^* + a_t)\omega^* - \frac{1}{2}a_t^2 + \varepsilon_t \quad (1)$$

Under this specification, the myopic optimal effort level is $a_t^* = \omega^*$. Nonetheless, because the agent does not know ω^* , he will choose $a_t = \mathbb{E}_t(\omega)$ where the expectation is taken with respect to the agent's belief at the beginning of period t . If he does not revise his effort choice for k periods, he will receive an average utility of $(\theta^* + a_t^*)\omega^* - \frac{1}{2}a_t^{*2}$ but he was expecting an average utility of $(\hat{\theta} + a_t^*)\omega^* - \frac{1}{2}a_t^{*2}$. In response, he will apply Bayes rule to update his beliefs about ω to get the posterior belief with $\mathbb{E}_{t+k}[\omega] = \frac{(\theta^* + \omega^*)\omega^*}{\hat{\theta} + \omega^*}$ which is lower than the initial belief. This will cause the agent to choose a lower effort at $t + k$. As a result, he will again receive an average utility that is lower than what he expected which will cause his belief to drift further down. This process will continue until the average utility equals his expected utility under the dogmatic belief that assigns probability 1 to $\hat{\theta}$. At that point, the student will have reached a self-defeating equilibrium and he will continue to choose sub-optimal effort forever.

³The example is illustrated for an overconfident agent but the results are symmetric for a digmatic agent who initially places probability one on some $\tilde{\theta} < \theta^*$.

Although the model of a dogmatic modeler rationalizes the prevalence of overconfident (underconfident) beliefs, the assumption that the agent has a degenerate belief and no mechanism through which he can update such belief is very restrictive. An alternative approach is proposed by Ba [2023]. She proposes an extension of the dogmatic agent who is able to jump from one dogmatic belief to another. By doing so, the agent might end up being dogmatic and correctly specified.

2.1.2 The Switcher

An agent is a *switcher* if they behave as a dogmatic, but is willing to entertain the possibility that they are of a different type. In particular, when they start off as a misspecified dogmatic, they are willing to switch to a different dogmatic belief if the data is convincing enough. Their prior is still degenerate and assigns probability one to a particular type, and zero to all other types. This means that a Bayesian update on θ does not change their beliefs about the type. However, they are willing to entertain two such beliefs and have a mechanism by which they decide which belief to adopt at any period t .

In order to abandon their initial dogmatic belief, the agent needs to observe a sequence of outcomes that are sufficiently unlikely to have happened if they were of the type they initially believed. In order to evaluate if the evidence is convincing enough, they keep track of the likelihood that each of the possible types generated the data. If the likelihood ratio is sufficiently large, the agent will switch to the alternative and behave as if they are dogmatic about the new type.

In particular, for an agent that starts with a dogmatic belief that they are of type $\hat{\theta}$ but is willing to consider the alternative explanation that they are of type $\tilde{\theta}$, the agent will switch to the alternative if:

$$\frac{p[h^t|\tilde{\theta}]}{p[h^t|\hat{\theta}]} > \alpha \geq 1$$

Where h^t is the history of outcomes up to time t and α is the switching threshold. By

keeping track of the likelihood ratio, the agent can perform a *Bayesian hypothesis test* and adopt the Dogmatic belief that best fits the data.⁴ Notice that if $\alpha \rightarrow \infty$, the behavior of the switcher will be indistinguishable from that of the Dogmatic modeler. In this sense, the switcher is a generalization of the dogmatic type.

By allowing the agent to keep track of the likelihoods and switching to an alternative type, the switcher can avoid the self-confirming equilibrium. However, if the prior belief on *omega* is sufficiently tight around a self-defeating equilibrium, the switcher might look identical to the dogmatic even in a case where α is not too large. This happens because under the agent’s prior, the likelihood ratio is unlikely to grow as fast as it is needed to escape the self-defeating equilibrium. In such situations, we say that the misspecified belief is persistent.

2.2 Framework 2

As in framework 1, the agent is of some type $\theta^* \in \Theta$ and the state is $\omega \sim F(\Omega)$. In this case, the agent chooses an action $a \in A$ and observes a binary outcome that is either a success or a failure. Denote the outcome by $o \in s, f$. The probability of observing a success is increasing in θ^* and in ω . Whenever the agent observes a success, he gets a payoff $v > 0$ and whenever the outcome is a failure, the payoff is 0. In addition, the probability of success is such that for each state, there is a unique optimal action that maximizes the agent’s expected payoff. Therefore, the probability of success can be seen as a monotone transformation of the utility from Framework 1.

I focus on two nested theories that have been widely studied within this framework: Full Bayesian updating and self-serving attribution bias. I explain each of these classical models of belief updating in what follows

⁴In a related problem @Schwarstein2021 proposes a similar updating procedure which relies on the Bayesian hypothesis test. However, in their model there is a sender who optimally chooses to propose a model that fits the data

2.2.1 The Bayesian

A Bayesian agent simultaneously updates their beliefs about θ and ω by using Bayes' rule. The posterior at period $t+1$ about θ after observing a history of outcomes h^t is given by:

$$p_{t+1}(\theta, \omega | h^t) = \frac{p[h^t | \theta, \omega] p_t(\theta, \omega)}{\sum_{(\theta', \omega')} p[h^t | \theta', \omega'] p_t(\theta', \omega')}$$

Where p_t is the belief at the start of period $t + 1$. The update is symmetric for ω .

Bayesian agents choose the effort level that maximizes their expected flow payoff by taking expectations over their prior beliefs about θ and ω . Since agents are myopic, even though all the parameters could be identified with enough variation in choices. Although the Bayesian agent is the closest to a fully rational agent discussed here, they might not learn their true type. This happens because, by being myopic, they do not internalize the tradeoff between flow payoff and learning. This can result in too little experimentation to learn their true type. An alternative to this approach is given by Hestermann and Yaouanq [2021] and is discussed in the concluding remarks.

Also notice that if a fully Bayesian agent has a dogmatic prior, they will never update their beliefs about the parameter that they are dogmatic about. This will imply that they are prone to the same types of errors as the dogmatic modeler in framework 1.

2.2.2 The Self-Serving Updater

A self-serving Bayesian is an agent who uses a biased version of Bayes rule to update their beliefs. They will update their beliefs about the state ω and his type θ simultaneously by over-attributing successes to a high value of θ and under-estimating the role of higher ω . Similarly, he will attribute failure to a low state to a greater degree than an unbiased agent would. To model the self-serving attribution bias, I take the approach of Benjamin [2019], where the posterior odds are given by:

$$p_{t+1}(\theta, \omega | h^t) = \frac{p[h^t | \theta, \omega]^{c(\theta, \omega, o_t)} p_t(\theta, \omega)}{\sum_{(\theta', \omega')} p[h^t | \theta', \omega']^{c(\theta', \omega', o_t)} p_t(\theta', \omega')}$$

with $c(\theta_H, \omega, o) < c(\theta_M, \omega, o) < c(\theta_L, \omega, o) \leq 1$ and $c(\theta, \omega_L, o) < c(\theta, \omega_M, o) < c(\theta, \omega_H, o) \leq 1$.

This formulation of the bias means that the agent will over-attribute successes to a high type and under-attribute them to a low type. Whereas they will over-attribute failures to a low state and under-attribute them to a high state. The parameter c determines the degree of bias and the restriction on the order ensures that the model is still falsifiable.

3 A Unifying Example

In order to compare the predictions of the theories discussed above, I develop a unifying example that allows me to isolate the forces behind each of the theories. The example is a modification of the one in Heidhues et al. [2018] and is adapted to be easier to implement in the lab.

The agent can be of one of 3 types: $\theta \in \{\theta_L, \theta_M, \theta_H\}$ with $\theta_H > \theta_M > \theta_L$. They face an unknown exogenous success rate $\omega \in \{\omega_L, \omega_M, \omega_H\}$ with $\omega_H > \omega_M > \omega_L$. Each of the values of ω is realized with equal probability. The agent knows the distribution of ω but not its realized value.

Denote the true type by θ^* and the true state by ω^* . The agent holds some prior belief about θ ⁵ and chooses a binary gamble $ein\{e_L, e_M, e_H\}$. The agent observes whether the gamble is a success or a failure and gets a payoff of 1 and if it is a success; they get 0 otherwise.

The probability of success is increasing in both θ and ω and is fully described by the following table:

Conditional on a type, the agent's flow payoff is maximized by choosing the gamble that

⁵which is potentially misspecified as in the dogmatic and switcher cases discussed above

	ω_H	ω_M	ω_L		ω_H	ω_M	ω_L		ω_H	ω_M	ω_L
e_H	50	20	2	e_H	80	50	5	e_H	98	65	25
e_M	45	30	7	e_M	69	65	30	e_M	80	69	35
e_L	40	25	20	e_L	65	45	40	e_L	75	55	45
	θ_L				θ_M				θ_H		

Table 1: Probability of success for each type, gamble and effort level

matches the state. For example, if the value of ω is ω_H , the agent's flow payoff is maximized by choosing e_H and if the state is ω_L the flow payoff is maximized by choosing gamble e_L , regardless of the value of θ . The agent myopically chooses gambles every period to maximize the flow payoff for $T < \infty$ periods.

After observing the outcome of each gamble, the agent updates their beliefs using some procedure and moves on to the next period.

Notice that both θ and ω can be identified from the outcomes if enough variation in the effort choices exists. This can be seen by confirming that there is no pair of θ and ω such that the probability of success is the same for all effort choices. Thus, by changing the effort choice, the agent can learn both their type and the state if they observe enough outcomes.

In this example, for an agent with a dogmatic belief about their type, a self-defeating equilibrium is one in which the agent chooses an effort level that, under the true θ , yields a frequency of success that is consistent with the agent's misspecified belief. That is $P[\text{sucess}|\theta^*, e^*] = P[\text{sucess}|\hat{\theta}, e^*]$ where e^* is the agent's myopic optimal choice.

In the data-generating process described above, there are five such equilibria. For example, if the agent is of type θ_M but mistakenly believes that he is of type $\hat{\theta} = \theta_H$ and the and $\omega^* = \omega_M$, when the effort chosen is e_L , the agent will observe a success with 45% chance. Because the agent dogmatically believes that their type is high, they will erroneously conclude that the rate is ω_L . Under this belief, the optimal action is e_L which will continue to generate successes with 45 probability, further reinforcing the incorrect belief. By doing so, the agent forgoes the payoff from gamble e_M which would yield a success with 65% chance.

By including self-confirming equilibria, the example captures the forces from each of the

updating mechanisms discussed in the previous section and allows for the comparison of the main forces behind the theories. For realizations of (θ, ω) for which there are self-confirming equilibria, the dogmatic agent will fall into the trap whereas the switcher will be able to escape it. Similarly, an agent with self-attribution bias will update their beliefs differently from an unbiased Bayesian, leading them to choose different gambles. I exploit such cases in order to test which model is a better fit for how subjects behave in a laboratory experiment.

⁶ In what follows I explain the details of how this example was implemented in the lab.

4 Experimental Design

I recruited undergraduate subjects from the CESS lab at NYU who participated in an in-person experiment. Sessions lasted approximately 45 minutes and subjects earned an average payment of \$22.6. The experiment was programmed using oTree [Chen et al., 2016].

The experiment consisted of 2 treatments: the *ego-relevant* condition and the *stereotype* condition. Subjects participated in only one of the treatments. All subjects within a session participated in the same treatment and the first 4 treatments were assigned the ego-relevant condition; the rest were assigned to the stereotype treatment. The tasks were identical across treatments, except for parameter θ . In the ego-relevant condition θ is the subject's own performance in a quiz, while in the stereotype condition, it is the performance of a randomly selected subject from another session.

The experiment had 3 parts. In Part 1 subjects had 2 minutes to answer as many multiple-choice questions as they could from a 20-question quiz. They did this for quizzes on 6 different topics. The topics were: Math, Verbal Reasoning, Pop-culture and Art, Science and Technology, US Geography, and Sports and Video Games. In this part, they did not know how many questions were available and they were given no feedback.

⁶because the setting does not match that of Heidhues et al. [2018], there will be situations for which the theory does not provide a prediction. If such cases arise in the lab, they will not be used for the analysis. However, whether a misspecified belief persists or not for the switcher, depends highly on the realized history of signals that he gets.

After taking all 6 quizzes, they proceeded to part 2 where they were asked to guess their score on each of them. In the stereotype treatment they were additionally asked to guess the score of a randomly drawn participant from a previous session. All they knew about the other participant was their gender identity and whether they were US nationals or not. For each guess, they had three score options: Low-Score (5 or fewer correct answers), Mid-Score (between 6 and 15 correct answers), High-Score (16 or more). Each of the score categories corresponded to θ_L , θ_M , and θ_H respectively. They were also asked to say how confident they felt about their choices. They had 4 possible answers: “it was a random guess”, “there is another equally likely score”, “I am pretty sure”, “I am completely sure”. These 4 answers are mapped to priors that place probabilities .33, .50, .75, and 1 to the chosen type. The remaining probability is split equally among the other two types. Questions in Part 2 were not incentivized, but subjects were told that providing an accurate answer would increase their chances of earning more money in the last part of the experiment.

The purpose of Part 2 is to classify subjects into overconfident, underconfident and correctly specified. If a subject guesses their score to be in a higher (lower) category than their true score, they are overconfident (underconfident); if they guess their score to be in the same category as their true score, they are correctly specified. This classification is done for each of the 6 topics separately.

Finally, in Part 3, subjects completed a belief updating task for each of the quizzes. Before starting the task they were reminded of their guess for the score. In the ego-relevant treatment, they were reminded of their guess about themselves and in the stereotype treatment they were reminded of their guess about the other participant. In the stereotype treatment, they were also reminded of the characteristics of the other participant.

For one topic at a time and in random order, they were presented with the three gambles from the example above and were asked to choose one of them. The probability of success was determined by their own score in the ego-relevant condition, and by the score of the other participant in the stereotype condition. Subjects had access to the three probability tables

in the printout of the instructions at all times and the meaning of each cell was explained in detail.

In the interphase, they had to choose which of the 3 tables they wanted to see before entering their choice in it. This was done as an alternative to a belief elicitation in each round. I take their choice of table to be indicative of their beliefs about the underlying type. I chose not to elicit the beliefs at each round to stay true to the forces in framework 1.

Once they have entered their choice, they observe a sample of 10 outcomes from the gamble they chose. After observing the outcomes, they returned to the choice screen and entered a new choice. In the choice screen subjects had access to the entire history of gambles and outcomes for that task as well as a summary of the outcomes so far. Once they entered 11 gambles (and observed 110 outcomes), they moved on to the next topic and repeated the same procedure. They all did this for all 6 topics.

At the end of the experiment, one of the 6 topics was randomly selected to determine the payment. They earned \$0.20 for each correct answer in the quiz, and for each success in the task in Part 3 for the selected topic.

Randomness is controlled throughout the experiment and sessions by setting a seed at the beginning of the first session. The seed was drawn at random and remained fixed for all sessions.⁷ By doing this I ensure that any two subjects who have the same type and face the same exogenous rate will observe the same outcomes and thus, if they use the same updating procedure, they should be choosing the same gambles. This design feature allows me to identify differences in updating procedures across subjects.

5 Predictions

In this section, I outline the behavior that is predicted by each of the theories discussed above. I start with the dogmatic modeler.

⁷The seed that was drawn at the beginning of session 1 was 3452. The same seed was used for all sessions. It is used both for drawing ω for each of the tasks in the experiment, as well as for drawing the outcomes from the gambles.

5.1 Dogmatic Modeler

Since the domain of the problem is discrete and finite in the example, the predictions of the original theory apply only to the combinations of parameters and initial beliefs for which there is a self-defeating equilibrium. In the example, there are 5 such combinations. For each of them, the dogmatic modeler predicts that the agent will fall into the self-defeating equilibrium and will be able to sustain the misspecified belief forever.

Table 2 describes the 5 self-defeating equilibria and the effort choices that sustain them. The first columns describe the combination of parameters and initial beliefs. The last column describes the effort choice that the agent will make in the long run. It is only for these combinations of parameter values and beliefs that the dogmatic model makes predictions within the Unifying example.

Prediction 1D: *Whenever an agent is of a type θ^* but mistakenly believes that they are of a type $\hat{\theta}$, and $(\theta^*, \omega^*, \hat{\theta})$ are such that there is a self-defeating equilibrium, the agent will fall into the trap and choose the effort level that sustains the misspecified belief forever.*

The model does not make predictions about what happens in cases where there is no stable belief. I assume that because there is no stable belief, the agent will eventually have to use some procedure to revise their belief about θ . In such cases, I aim to determine which of the alternative explanations provided by the other theories is a better fit for the data.

True Type (θ^*)	True State (ω^*)	Believed Type ($\hat{\theta}$)	Believed state ($\hat{\omega}$)	Effort
θ_L	ω_H	θ_M	ω_L	e_L
θ_M	ω_L	θ_L	ω_M	e_M
θ_M	ω_M	θ_H	ω_L	e_L
θ_M	ω_M	θ_L	ω_H	e_H
θ_M	ω_H	θ_H	ω_M	e_M

Table 2: Stable beliefs and the effort choices that support them for the unifying example

Although the dogmatic model does not apply to all possible parametrizations and beliefs, whether subjects fall into the self-defeating equilibria or not is still informative of the updating procedure that they are using. Understanding if the presence of traps is a key feature

preventing subjects from learning the optimal action is important for understanding the prevalence of overconfidence. Similarly, gaining insight into what happens when there are no traps is important for understanding what are the other reasons why overconfidence might arise and prevail.

5.2 Switcher

Since the switcher starts as the same dogmatic agent, the initial behavior of both types of agents is identical, however, because the switcher is keeping track of the likelihood ratio, they will be able to escape the self-defeating equilibrium if the evidence is convincing enough. Therefore there is a positive probability that the switcher will adjust their initially misspecified belief about θ and learn the true state.

Prediction 1S: With positive probability, the switcher will escape the self-defeating equilibrium and learn the true state.

One caveat is that when the switcher and the dogmatic agent both start with a correctly specified belief, neither of them will fall into the self-defeating equilibria and thus will look identical even in the long run. This means that in order to distinguish between the two theories, I need to look at cases where the agent starts with a misspecified belief.

The probability that the switcher will escape the self-defeating equilibrium depends on the prior belief about ω . If the prior is sufficiently tight around the self-defeating equilibrium, the likelihood ratio will not grow fast enough, however, in the example above, the prior is uniform over the states and therefore, the likelihood ratio is more likely to grow fast enough.

5.3 Self-Attribution Bias

The key feature of Self-attribution bias is the asymmetric treatment of good news and bad news. In particular, the agent will over-attribute successes to a high type and under-attribute them to a low type. Similarly, they will over-attribute failures to a low state and under-attribute them to a high state. This implies that, after observing a failure, the agent will

adjust their effort downwards by more than what an unbiased Bayesian would have done. In contrast, after observing a success, the agent will adjust their effort upwards by less than a Bayesian would have done and if the bias is large enough, it could be that the agent will not adjust their effort upwards at all, or even decrease their effort in response to a success.

This is in stark contrast with what a Dogmatic modeler would do. A dogmatic modeler always attributes any variation in the outcome to the state and never updates their beliefs about θ . Therefore, they will always increase their effort choices after a surprising success and decrease it after a failure.

On the other hand, the biased behavior can be in line with the behavior of a switcher. In particular, if the agent starts with a misspecified belief on $\hat{\theta}$, and is willing to switch to a belief with $\theta' > \hat{\theta}$, whenever the paradigm shift happens, it will likely be in response to a surprising streak of successes. In this case, when they adjust their belief about θ , they will also adjust their effort choice. Since they were initially underconfident, they had been choosing an effort that was too high relative to the true state, and therefore, the effort is likely to fall in response to a surprising streak of success.

Prediction 1A: *After observing a failure, the agent will adjust their effort downwards by more than a Bayesian would have done. After observing a success, the agent will adjust their effort upwards by less than a Bayesian would have done. If the bias is large enough, the agent might not adjust their effort upwards at all, or even decrease their effort in response to a success.*

Prediction 1A helps distinguish between the self-attribution bias and the dogmatic modeler. However, it does not help distinguish between the self-attribution bias and the switcher. In order to do so, I need to look at cases where the agent starts with a correctly specified belief. In these cases, the switcher is very unlikely to switch to a different belief and therefore, the behavior of the switcher will mostly change in the same direction as the information he receives.⁸ Instead, the self-attribution bias will lead to a different behavior. In particular,

⁸If they get a streak of failures they are likely to adjust their effort downwards and if they get a streak of successes they are likely to adjust their effort upwards.

even for agents who have a belief that is correctly specified and place a lot of weight on the correct values of θ , the bias will cause these agents to become overconfident after a sequence of successes. This force is not present in the switcher or the dogmatic and it is the key feature that allows me to identify a biased updating procedure.

Prediction 2A: *After observing a streak of successes, the agent will update their belief about θ upwards. This will lead to the possibility of observing subjects who start correctly specified and become overconfident.*

5.4 The Bayesian

The Bayesian serves as the closest benchmark to a fully rational agent. However, as discussed above, they are still vulnerable to falling into self-defeating equilibria, in which case, their long-run behavior will be identical to that of the dogmatic modeler. The main distinguishing feature is that when they do fall into a self-defeating equilibrium, there is no bias in the direction of the misspecification that they end up being trapped in. In contrast, the dogmatic modeler will always be trapped in a misspecified belief that they started with and the self-attribution updater will always become overconfident.

Prediction 1B: *When an agent starts with a diffused prior and falls into a self-defeating equilibrium, they will be equally likely to fall into the underconfident trap as the overconfident trap.*

Prediction 2B: *Bayesian agents always update their effort choices in the same direction as the information they receive. That is, if they see successes, they would only increase their effort; and if they see failures they would only decrease their effort.*

6 Stereotypes

So far I have focused only on the case of overconfidence. However, each of the models can also be used to explain the prevalence of underconfidence. An agent can have a dogmatic

belief about the ability of a particular group of people and either under or overestimate the associated parameter. Similarly, they can be willing to switch between two dogmatic beliefs about their ability as a switcher would. Finally, they can have a biased updating procedure in which they treat good news and bad news asymmetrically.

As such, all the predictions discussed above apply to the case of Stereotypes as well. My analysis of stereotypes will focus only on the degree to which ego-relevance of the type affects the updating procedure and not on the motivations behind the bias. I am particularly interested in whether the explanatory power of the models differs with the degree of ego-relevance of the type.

7 Results

7.1 Initial Beliefs and misspecifications

As mentioned in the predictions section, the dogmatic modeler and the switcher are only distinguishable when the agent starts with a misspecified belief. Figure 1 shows the distribution of misspecifications by treatment. The histogram considers the difference between the subject’s true score in the quiz and their guess about their score. If their guess is in a higher category than their true score, they are overconfident; if it is in a lower category, they are underconfident; and if it is in the same category, they are correctly specified. Overall, 0.43 of the guesses made in the questionnaire were misspecified. The rest of the guesses coincided with the true score. And the distribution of misspecifications is similar across treatments.

Although the overall distribution is similar for both treatments, the misspecifications arise for different combinations of characteristics in each treatment. Figure 2 shows a heatmap of the misspecifications that arose in each treatment. In the stereotype treatment, subjects are most underconfident about the performance of other non-American females in Sports and Video Games. In contrast, they are most overconfident about the performance of American men in Verbal Reasoning. In the ego-relevant treatment, subjects are most underconfident

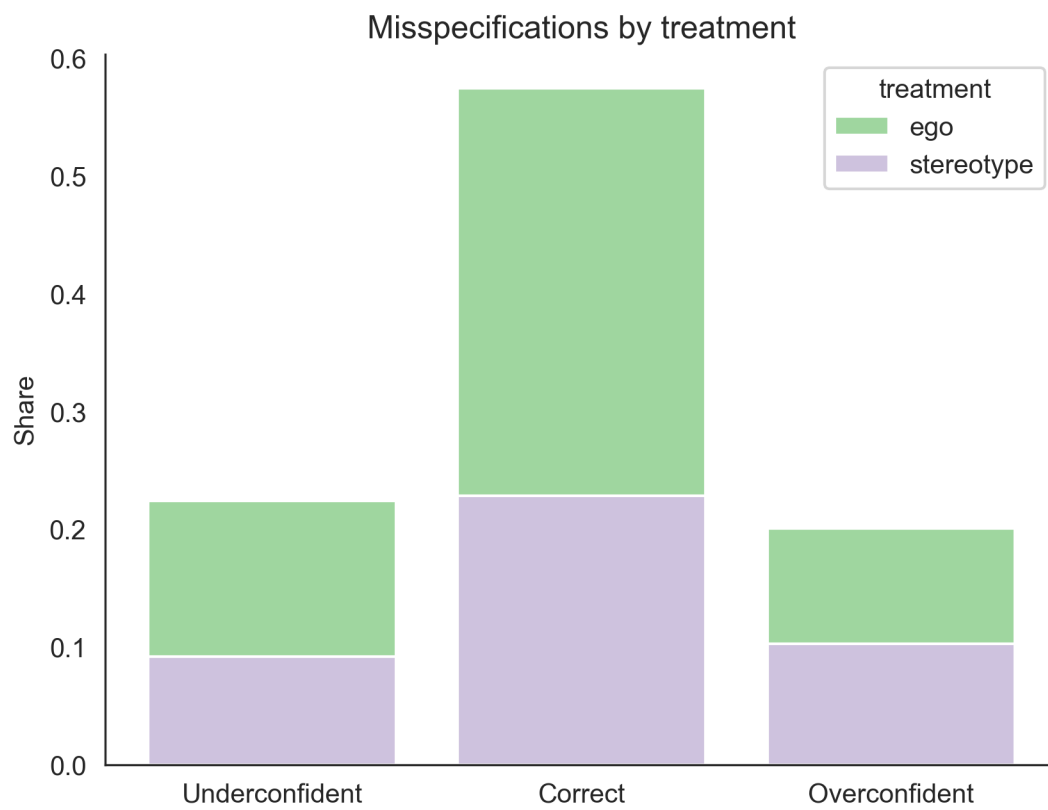


Figure 1: Initial misspecifications by treatment

about their own performance in pop culture and art, while most overconfident about their performance in Verbal Reasoning.




Figure 2: Initial misspecifications by treatment and characteristics

Having widespread misspecifications in both treatments is important for the analysis because it allows me to distinguish between the dogmatic modeler and the switcher. If there were no misspecifications, the two models would be indistinguishable.

7.2 Learning

In this section, I analyze the learning behavior of subjects in the experiment. There are two parameters that subjects can be learning about: the exogenous parameter ω and the type θ . Their belief about the exogenous parameter is tracked by their choice of effort, while their belief about the type is tracked by their choice of a matrix in which to enter their effort.

7.2.1 Learning about the state

In order to analyze the learning about the state, I look at the share of optimal choices that are made at each round. I find that although subjects seem to be improving in their choices overall, the last choice coincided with the true state only in 52% of the choices in the last round of each task. This is statistically greater than the initial share of optimal choices

($p < 0.01$). However, it is still far from complete learning. It is also important to note that learning is similar across treatments. The share of choices that were consistent with the true state for each round is reported in Figure 3.




Figure 3: Share of optimal choices by round

A closer look at whether people learned or not reveals that there is a good amount of heterogeneity in the sample. Figure 4 shows the share of optimal choices by round for subjects that chose an effort that matched the state in 3 out of the last 4 rounds. It also shows the share of optimal choices for subjects who chose an effort that matched the state in fewer than 3 out of the last 4 rounds. I label the former as learners and the latter as non-learners, with learners making up 39% of the sample.

In what follows I will try to disentangle the reasons for the lack of learning. I will argue that it is not due to the presence of traps, but rather due to biased updating. According to




Figure 4: Share of optimal choices by round for subjects who learned (their effort in at least 3 out of the last 4 choices matched the state), and subjects who did not learn

the theories, the main reason why subjects might not learn is that they have an incorrect belief about the type or they develop an incorrect belief about the type.

7.2.2 Learning about the type

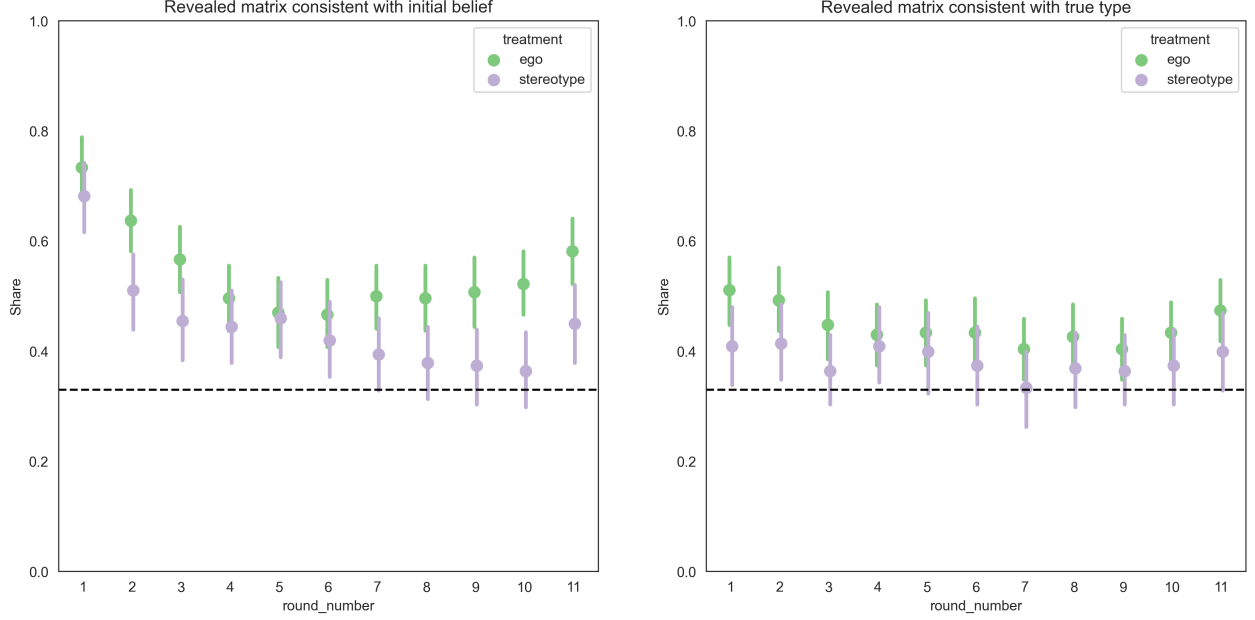


Figure 5: Matrix choices by round

Since the elicitation of the belief about the type was not incentivized and not elicited in a standard way, I first need to confirm that subjects were not just randomly choosing a matrix in which to enter their effort. The left panel of Figure 5 shows the share of subjects who chose a matrix consistent with their initial reported belief. In round 1, 71% of the subjects chose a matrix that was consistent with their initial belief. This indicates that subjects were not just randomly choosing a matrix in which to enter their effort. From round 2 onwards, the share steadily declines, but still not as far as to indicate a random choice of matrices. This is consistent with the subjects moving away from their starting belief through some updating procedure.

On the right panel of Figure 5 I show the share of subjects who chose a matrix that is consistent with their true type. Unlike the left panel, there is no clear trend, which indicates

that although they are moving away from their initial belief, they are not moving towards their true type, which means that overall, misspecification is not decreasing. A closer look at the data reveals a good amount of heterogeneity in the underlying behavior.

Figure 6 shows the transition matrix for subjects who started in each of the 3 possible starting specifications and the specifications that they ended up in at the end of the updating task. The two things to note are that the initial belief is the most likely end belief. This is consistent with some degree of stickiness of the misspecifications as the switcher and dogmatic models would predict. However, the data also presents a lot of subjects who started with a correct belief of the type and ended up overestimating it. This is consistent only with the self-attribution bias. Lastly, the subjects who initially overestimate the score, are the least likely to learn their true type.

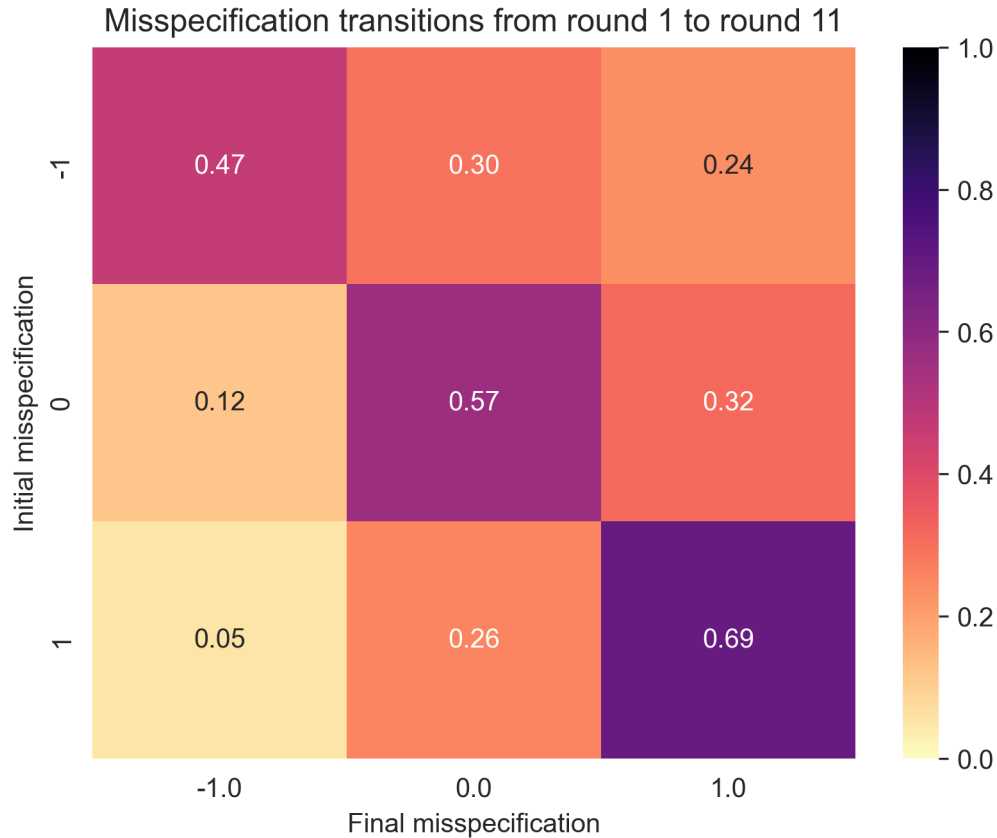


Figure 6: Transition matrix for subjects who started in each of the 3 possible starting specifications and the specification that they ended up in at the end of the updating task

The transition matrix points towards the presence of self-attribution bias in the data. However, it is still possible that there is heterogeneity in the updating procedure across subjects. In the next section I will look at the role of self-defeating equilibria in preventing learning. I find that the presence of self-defeating equilibria does not seem to be the main deterrent to learning.

7.3 Learning traps

As described above, the presence of stable misspecified beliefs act as learning traps. If an agent falls into one of these traps, bayesian updating will lead them to choose the same effort forever. These self-defeating equilibria allow me to distinguish between the dogmatic modeler and the switcher. If the agent falls into a trap and stays there forever, they are more likely to be a dogmatic modeler. If they escape the trap, they are more likely to be a switcher or to be updating with some attribution bias.

Overall, it does not seem to be the presence of traps that is preventing learning. Figure 7 shows that the share of subjects who learned in cases where there was a trap is larger than the share of subjects who learned in cases where there was no trap. This is true for both treatments. This indicates that even in the presence of traps, some people learn.

A closer look at the behavior of only the cases in which there were traps reveals that among those who did not learn, the choices that support stable beliefs were chosen much more often than among subjects who learned. Figure 8 shows the share of choices that were consistent with a stable belief for the learners and the non-learners. The difference is statistically significant ($p < 0.01$). Thus, although the presence of self-defeating equilibria is not the main deterrent to learning, it does seem to be a factor in the behavior of those who do not learn.

When looking at the behavior in the last 4 rounds of each task, only 16% of the subjects were trapped⁹ in a self-defeating equilibrium. These subjects could have been dogmatic,

⁹I consider a subject to be trapped if 3 out of the last 4 choices were consistent with the self-defeating equilibrium




Figure 7: Share of subjects who learned in cases where there was a trap and in cases where there was no trap

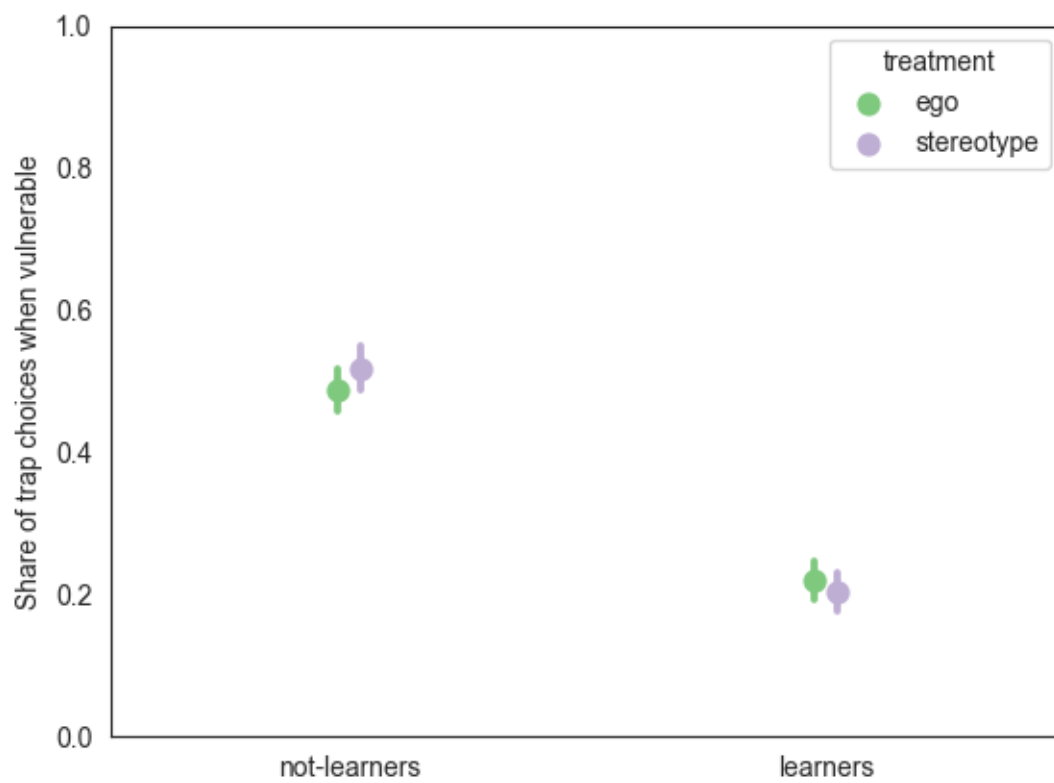


Figure 8: Share of choices that were consistent with a stable belief for the learners and the non-learners

switchers for whom the likelihood did not grow fast enough, or Bayesian agents who fell into a self-defeating equilibrium. The rest of the subjects were either updating their beliefs with a bias or were switchers.

So far I have accounted for 16% subjects as being trapped, and 39% having learned the state correctly. From the remaining 45%, 60% were prone to traps but did not fall into them, therefore they could not have been dogmatic modelers. Alternative explanations are that they were switchers who considered an incorrect alternative model, or that they were updating with a bias so large that their effort choices were far from optimal.

In the following section, I look at indications of biased updating among all the groups identified so far.

7.4 Reactions to good and bad news

One of the key features of the self-attribution bias is that the agent will react differently to good news and bad news. In particular, they will overreact to bad news and underreact to good news. I define good news to be any sample of outcomes for which the realized number of successes is greater than the average number of successes so far. Similarly, I define bad news to be any sample of outcomes for which the realized number of successes is smaller than the average number of successes so far. In addition, I calculate the difference between the realized number of successes and the average number of successes so far and call it the news differential. Figure 9 shows the correlation between change in effort and the news differential, separately for good and bad news.

The overall pattern is that bad news are associated with a decrease in effort, and a more negative news differential is associated with a larger decrease in effort. On the contrary, good news are associated with a small decrease in effort. As mentioned in the predictions section, this is consistent with the self-attribution bias as well as with a change of paradigm. It is not consistent with the dogmatic modeler or the Bayesian. Because I do not directly elicit beliefs, it is not possible to distinguish between the self-attribution bias and the change

of paradigm with this data alone. Nevertheless, the fact that about a third of the subjects started with a correctly specified belief and ended up overconfident is consistent only with the self-attribution bias. Therefore, I conclude that the self-attribution bias is the most likely explanation for the behavior of the subjects in the experiment.




Figure 9: Change in effort by news differential

A regression analysis of this same data reveals that indeed, the slopes are different for the two types of news and that the difference is statistically significant ($p < 0.01$). The regression results are reported in Table 3.

Table 3: Regression of effort change on news differential

	<i>Dependent variable:</i>		
	All	Change in effort Ego-relevant	Stereotype
	(1)	(2)	(3)
Good News	−0.12** (0.05)	−0.16*** (0.05)	−0.05 (0.05)
News differential	0.22*** (0.02)	0.22*** (0.02)	0.21*** (0.02)
News Differential * Good News	−0.27*** (0.02)	−0.25*** (0.02)	−0.29*** (0.02)
Constant	0.31*** (0.04)	0.31*** (0.04)	0.30*** (0.04)
Observations	4,680	2,700	1,980
R ²	0.04	0.04	0.04
Adjusted R ²	0.04	0.04	0.04
<i>Note:</i> *p<0.1; **p<0.05; ***p<0.01			

8 Structural estimation

8.1 Model fit

9 Conclusion

I designed a setting in which I can distinguish between the forces of dogmatism, paradigm shifts, and optimal expectations. Through a laboratory experiment, I collected data on behavior and infer beliefs from it. I find that the behavior of the subjects is consistent with the self-attribution bias. In particular, I find that subjects decrease their effort in response to bad news as well as to bad news, although the latter decrease is smaller. This is consistent with the self-attribution bias as well as with a change of paradigm. However, I also observe that about a third of the subjects started with a correctly specified belief and ended up overconfident. This is consistent only with the self-attribution bias. Therefore, I conclude that the model that best explains the data is that of optimal expectations.

I have small or no effects throughout on the ego-relevance of the type. Although the motivation behind the theories does not directly apply to the case of stereotypes, the lack of stark differences suggest that biases arise due to an incorrect updating procedure and not necessarily due to the self-serving motivation behind the bias. The lack of differences in the behavior of the subjects across treatments could also be due to self-censoring of the subjects. In particular, subjects might be initially reluctant to express their true belief about a particular group of people. In addition, as they learn and realize that their beliefs might be incorrect, they might overreact to signals to try to compensate for the initial bias. Such an updating procedure would be consistent with the self-attribution bias for the case of stereotypes. I remain agnostic about whether that is truly the case or whether the lack of differences is due to other factors.

Another bias that has been studied and suggested in the literature is the confirmation bias. In particular, it has been suggested that subjects who are overestimate θ in their

prior, will over-attribute successes to a high type and under-attribute them to a low type. Symmetrically, subjects who initially underestimate θ will over-attribute failures to a low type and under-attribute them to a high type. I do not find evidence of such a bias in the data. This is consistent with what ? find in their experiment.

I further investigate possible heterogeneity in the underlying models being used by estimating the structural parameters of the models.

References

Cuimin Ba. Robust misspecified models and paradigm shifts. 2023.

Daniel J. Benjamin. *Errors in probabilistic reasoning and judgment biases*, pages 69–186. 2019. doi: 10.1016/bs.hesbe.2018.11.002.

Daniel L. Chen, Martin Schonger, and Chris Wickens. otree—an open-source platform for laboratory, online, and field experiments. *Journal of Behavioral and Experimental Finance*, 9:88–97, 3 2016. ISSN 2214-6350. doi: 10.1016/J.JBEF.2015.12.001.

Paul Heidhues, Botond Köszegi, and Philipp Strack. Unrealistic expectations and misguided learning. *Econometrica*, 86:1159–1214, 2018. ISSN 0012-9682. doi: 10.3982/ecta14084.

Nina Hestermann and Yves Le Yaouanq. Experimentation with self-serving attribution biases. *American Economic Journal: Microeconomics*, 13:198–237, 2021. ISSN 19457685. doi: 10.1257/mic.20180326.