

Learning with Misspecified Models: the case of overconfidence

Jimena Galindo

October 2, 2023

OVERCONFIDENCE: Belief that my type is higher than it truly is (“overestimation” as in Moore and Healy (2008))

Overconfidence is Costly

OVERCONFIDENCE: Belief that my type is higher than it truly is (“overestimation” as in Moore and Healy (2008))

Seems to be persistent in various settings.

- Excess entry of entrepreneurs (Camerer and Lovo, 1999)
- Suboptimal genetic testing and savings (Oster et al. 2013)
- Workers overestimate their productivity (Hoffman and Burks, 2020)

Ultimately it leads to sub-optimal choices

Models of Learning

Focus on setting with 2 parameters:

- An **Ego-Relevant** parameter
- An **Exogenous** parameter

Some of the features that theory has incorporated to explain overconfidence are:

- Dogmatism
- Paradigm shifts
- Motivated beliefs
- Myopic optimization

Four Theories of Misspecified Learning

1. **Self-defeating equilibrium** (Heidhues et al. (2018))
 - Bayesian about exogenous parameters
 - Dogmatic about ego-relevant parameters
2. **Bayesian hypothesis testing** (Schwarstein and Sunderam (2021), Ba (2022))
 - Bayesian about exogenous parameters
 - Paradigm shift for ego-relevant parameters
3. **Motivated Beliefs / Self-Attribution Bias** (Brunnermeier and Parker (2005), Bracha and Brown (2012))
 - Optimally biased updating
 - Utility from held beliefs
4. **Myopic Bayesian** (Hestermann and Le Yaouanq, (2021))
 - Bayesian about both
 - Maximizes flow utility only

Which of the proposed features better explain the observed behavior?

- Do we observe heterogeneity in the use of misspecified models?

Is ego-relevance of the parameter a key feature for the misspecification?

- Are ego-relevant misspecifications more likely to persist than stereotypes?
- Can the same theories be used to explain the prevalence of stereotypes?

An Example (from Heidhues et al. (2018))

A student has unknown **intrinsic ability** θ^* (ego-relevant parameter)

They choose a level of **effort** $e \geq 0$ (choice)

Effort and ability are evaluated by a **grading system** ω (exogenous parameter)

The student wants to maximize:

$$u(e) = (\theta^* + e)\omega - \frac{1}{2}e^2 + \varepsilon$$

Regardless of their own type and of their beliefs about it, they should choose

$$e^*(\omega) = \omega$$

Learning is Possible

This exercise is repeated for $t = 0, 1, \dots$

$$y_t = (\theta^* + e_t)\omega - \frac{1}{2}e_t^2 + \varepsilon_t$$

Note that both parameters are identified in this setting:

- Choosing \hat{e} and $\hat{e} + 1$ over multiple periods allows identification of ω
- Once ω is known, θ can be backed out

How come people don't learn their true type and don't choose the optimal effort?

Road-map

1. Unifying Framework
2. Mechanisms and Predictions
3. Experimental Design
4. The Data
5. Parameter Estimation
6. Results

Framework

A Unifying Framework

Ego-relevant parameter: $\theta \in \{\theta_H, \theta_M, \theta_L\}$

Exogenous parameter: $\omega \in \{\omega_H, \omega_M, \omega_L\}$ with $p(\omega_k) = 1/3$

Choices: $e \in \{e_H, e_M, e_L\}$

Binary Outcomes: $s_t \in \{\text{success}, \text{failure}\}$ with $p[\text{success}|e, \omega, \theta]$ and p is an order-preserving transformation of $u(x)$

The Data Generating Process

The probability of success is given by:

	ω_H	ω_M	ω_L
e_H	50	20	2
e_M	45	30	7
e_L	40	25	20
	θ_L		

	ω_H	ω_M	ω_L
e_H	80	50	5
e_M	69	65	30
e_L	65	45	40
	θ_M		

	ω_H	ω_M	ω_L
e_H	98	65	25
e_M	80	69	35
e_L	75	55	45
	θ_H		

The Data Generating Process

	ω_H	ω_M	ω_L
e_H	50	20	2
e_M	45	30	7
e_L	40	25	20
	θ_L		

	ω_H	ω_M	ω_L
e_H	80	50	5
e_M	69	65	30
e_L	65	45	40
	θ_M		

	ω_H	ω_M	ω_L
e_H	98	65	25
e_M	80	69	35
e_L	75	55	45
	θ_H		

The Data Generating Process

Diagram illustrating the Data Generating Process, showing three payoff matrices for players e_H , e_M , and e_L across different states of nature (θ_L , θ_M , θ_H).

Each matrix has columns labeled ω_H , ω_M , and ω_L , and rows labeled e_H , e_M , and e_L . The values in the cells represent payoffs.

Payoff Matrix for θ_L :

	ω_H	ω_M	ω_L
e_H	50	20	2
e_M	45	30	7
e_L	40	25	20

Payoff Matrix for θ_M :

	ω_H	ω_M	ω_L
e_H	80	50	5
e_M	69	65	30
e_L	65	45	40

Payoff Matrix for θ_H :

	ω_H	ω_M	ω_L
e_H	98	65	25
e_M	80	69	35
e_L	75	55	45

Arrows indicate the flow of information from the state of nature to the players' payoffs.

A Stable Misspecified Belief

	ω_H	ω_M	ω_L
e_H	50	20	2
e_M	45	30	7
e_L	40	25	20
	θ_L		

	ω_H	ω_M	ω_L
e_H	80	50	5
e_M	69	65	30
e_L	65	45	40
	θ_M		

	ω_H	ω_M	ω_L
e_H	98	65	25
e_M	80	69	35
e_L	75	55	45
	θ_H		

The Stable Beliefs

	ω_H	ω_M	ω_L
e_H	50	20	2
e_M	45	30	7
e_L	40	25	20

θ_L

	ω_H	ω_M	ω_L
e_H	80	50	5
e_M	69	65	30
e_L	65	45	40

θ_M

	ω_H	ω_M	ω_L
e_H	98	65	25
e_M	80	69	35
e_L	75	55	45

θ_H

Mechanisms and Predictions

An Example

- True type is θ_M
- True parameter is $\omega_M \rightarrow$ the student believes it is uniformly distributed

	ω_H	ω_M	ω_L
e_H	50	20	2
e_M	45	30	7
e_L	40	25	20
	θ_L		

	ω_H	ω_M	ω_L
e_H	80	50	5
e_M	69	65	30
e_L	65	45	40
	θ_M		

	ω_H	ω_M	ω_L
e_H	98	65	25
e_M	80	69	35
e_L	75	55	45
	θ_H		

The Dogmatic Modeler

Holds a degenerate belief: type is $\hat{\theta}$ with probability 1

Their belief is potentially misspecified:

- Overconfident if $\hat{\theta} > \theta^*$
- Underconfident if $\hat{\theta} < \theta^*$

Updates $p_t(\omega)$ using Bayes Rule

$$p_{t+1}(\omega|s, \hat{\theta}) = \frac{p_t(s_t|\omega, \hat{\theta})p_t(\omega)}{\sum_{\omega'} p_t(s_t|\omega', \hat{\theta})p_t(\omega')}$$

The Dogmatic Modeler: Mechanism

A student who dogmatically believes he is θ_H but truly is θ_M

The exogenous parameter is ω_M

1. Chooses e_H and is disappointed \rightarrow adjust belief about ω downward
2. Eventually chooses e_M and is disappointed as well \rightarrow adjust belief about ω
3. Eventually chooses e_L and falls into a self-confirming equilibrium

	ω_H	ω_M	ω_L
e_H	50	20	2
e_M	45	30	7
e_L	40	25	20
	θ_L		

	ω_H	ω_M	ω_L
e_H	80	50	5
e_M	69	65	30
e_L	65	45	40
	θ_M		

	ω_H	ω_M	ω_L
e_H	98	65	25
e_M	80	69	35
e_L	75	55	45
	θ_H		

Dogmatic Overconfident: Simulated




Figure 1: $\theta^* = \theta_M$, $\hat{\theta} = \theta_H$, $\omega^* = \omega_M$

The Switcher (paradigm shifts)

Same initial belief as the Dogmatic, but is willing to consider an alternative paradigm θ'

Keeps track of the likelihoods of the two possible paradigms:

- $p_t(s_t|\cdot)$ for $\hat{\theta}$ and θ'

They switch to whichever paradigm is more likely to have generated the signals

$$\frac{p_t(s_t|\theta')}{p_t(s_t|\hat{\theta})} > \alpha \geq 1$$

The Switcher: Mechanism

1. Chooses e_H and is disappointed \rightarrow adjust belief about ω downward
2. Eventually chooses e_M and is disappointed as well \rightarrow adjust belief about ω
3. Avoids the self-defeating equilibrium if the likelihood of θ_M becomes larger than that of θ_H

Switcher Overconfident: Simulation

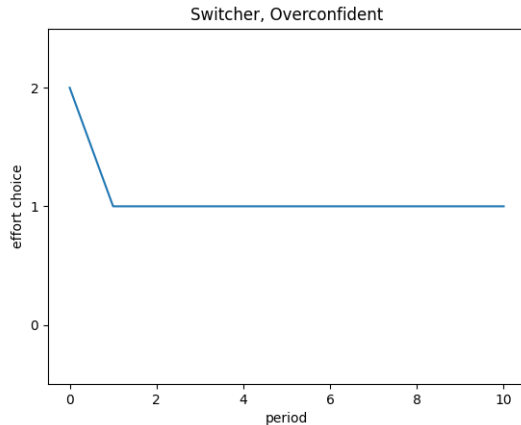


Figure 2: $\theta^* = \theta_M$, $\hat{\theta} = \theta_H$, $\omega^* = \omega_M$, $\alpha = 1.1$

Self-Attribution Bias / Optimal Expectations

Start with a diffused prior over (θ, ω) but updates with a bias

$$p_{t+1}(\theta, \omega | s_t) = \frac{p_t(s_t | \theta, \omega)^{c(\theta, \omega, s_t)} p_t(\theta, \omega)}{\sum_{(\theta', \omega')} p_t(s_t | \theta', \omega')^{c(\theta', \omega', s_t)} p_t(\theta', \omega')}$$

Bias is such that


$$c(\theta_H, \omega, \text{good news}) \leq c(\theta_M, \omega, \text{good news}) \leq c(\theta_L, \omega, \text{good news}) \leq 1 \quad \forall \omega$$

And

$$c(\theta, \omega_L, \text{bad news}) \leq c(\theta, \omega_M, \text{bad news}) \leq c(\theta, \omega_H, \text{bad news}) \leq 1 \quad \forall \theta$$

1. Chooses e that maximizes utility according to priors
 - Belief on $\mathbb{E}[\omega]$ deteriorates a lot after bad news \rightarrow big change in effort
 - Belief on $\mathbb{E}[\theta]$ increases a lot after good news \rightarrow small positive (or negative) change in effort

Self-Attribution: Simulation




Start with a diffused prior over (θ, ω) and updates correctly

$$p_{t+1}(\theta, \omega | s_t) = \frac{p_t(s_t | \theta, \omega) p_t(\theta, \omega)}{\sum_{(\theta', \omega')} p_t(s_t | \theta', \omega') p_t(\theta', \omega')}$$

But if they start with a prior that is “tight” around a self-defeating equilibrium they will never learn

All Models

Mid Type, rate = 1



Experimental Design

Set the Types

- Quiz: Answer as many questions as you can in 2 minutes
 - Math, Verbal, Pop-Culture, Science, Us Geography, Sports and Video games
- How many questions do you think you answered correctly in each quiz?
 - 0 to 5 (θ_L)
 - 6 to 15 (θ_M)
 - 16 or more (θ_H)
- How sure are you about your guess?
 - Random guess $\rightarrow 1/3$
 - Another is equally likely $\rightarrow 1/2$
 - Fairly certain $\rightarrow 3/4$
 - Completely sure $\rightarrow 1$

Effort choice and feedback (One topic at a time)

- Choose a gamble: A, B or C (effort)
- Receive a sample of 10 signal realizations

x 11 per topic

Stereotype condition

Observe the characteristics of a participant

- Gender,
- US National or not

Answer the same questions about self and other

Belief updating and effort choice:

- The DGP depends on the θ the other participant

x 11 per topic

Eliciting Beliefs?

- $E[\omega]$ is revealed by their choice of effort
- Eliciting beliefs for θ can incentivize learning in a way that is not consistent with the theory

Allow them to see the success rate matrix for only one type.

- Track the matrices they choose to see in each round

Based on the other participant's Science and Technology Quiz results

Which probability matrix would you like to see?

Low Score

Mid Score

High Score

Your Previous Outcomes

Choice

Successes

Failures

You have no data for this task yet

See History

Next

Based on the other participant's Science and Technology Quiz results

Which probability matrix would you like to see?

Low Score

Mid Score

High Score

Choose a gamble :		Rate A	Rate B	Rate C
A	<input type="radio"/>	40	45	65
B	<input type="radio"/>	30	65	69
C	<input type="radio"/>	5	50	80

Your Previous Outcomes

Choice

Successes

Failures

You have no data for this task yet

See History

Next

The Data

The Data

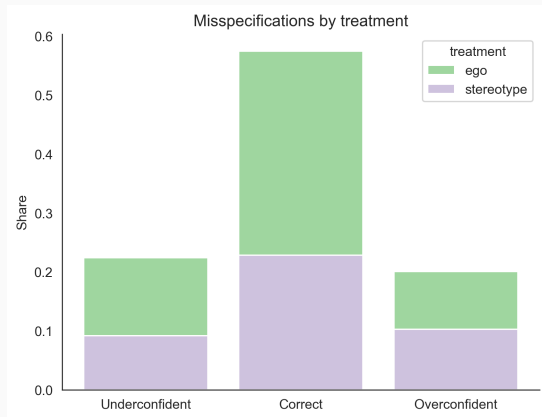
Subject pool:

- Run at the CESS lab in person
- 45 subjects in Ego
- 33 subjects in Stereotype

The Sessions:


- 8 sessions
- 45 minutes on average
- Average payment: \$23
 - \$10 show-up fee
 - \$0.20 per correct answer
 - \$0.20 per success
 - Paid one topic at random

Initial Misspecifications




certainties


The Stereotypes




misspecifications



Learning \ominus

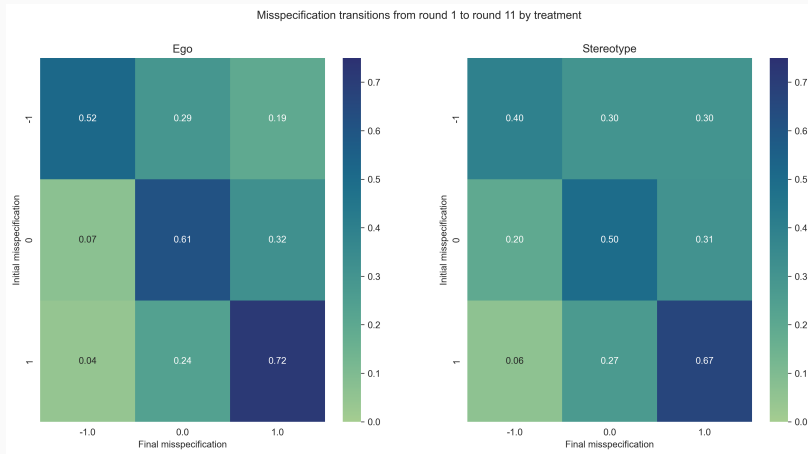


Changes in Misspecifications




overall

Transitions



Good News v. Bad News



Parameters

Identification of α

Whenever the agent switches from one paradigm to another, they are revealing that

$$\frac{p_t(s^t|\theta')}{p_t(s^t|\hat{\theta})} = \alpha$$

Notice that this identifies an upper bound for α

I take the average value of the likelihood ratio when the agent changes their choice of θ to be α

I find $\alpha = 1.48$ and no difference across treatments

Calibration of Bias

Simulation on a grid of parameters

For each task take the parameters that minimize the distance between the simulated and the actual effort

Average for each subject

Average across subjects

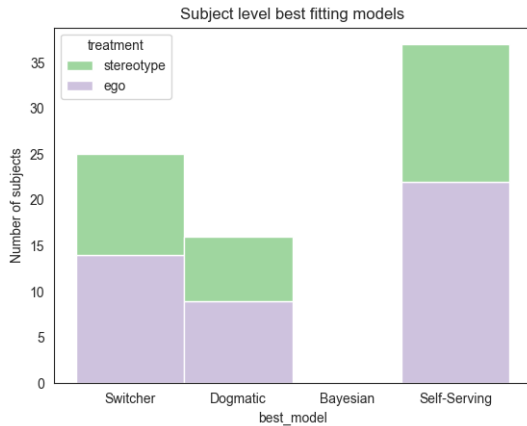
$$c(\theta_H, \omega, \text{good news}) = c(\theta, \omega_L, \text{bad news}) = 0.137$$

$$c(\theta_M, \omega, \text{good news}) = c(\theta, \omega_M, \text{bad news}) = 0.36$$


$$c(\theta_L, \omega, \text{good news}) = c(\theta, \omega_H, \text{bad news}) = 1$$

Heterogeneity

Model Fit: Distributions



Model Fit: Distance



Concluding Remarks

- Data is not fully consistent with the models
- Some indications of self-attribution bias and/or paradigm shifts
- Need a better estimation of the parameters
- Are subjects experimenting in order to learn?
 - Hestermann and Le Yaouanq (2021)

What is Next


1. Have a better estimation of the attribution bias parameters
 - Estimate using SMM
 - Elicit beliefs within this framework
2. Can dynamic learning explain the data better?
 - This model would predict underconfidence to be more persistent than overconfidence

The end


Thank you!

Misspecifications

misspecifications by topic and characteristics before feedback

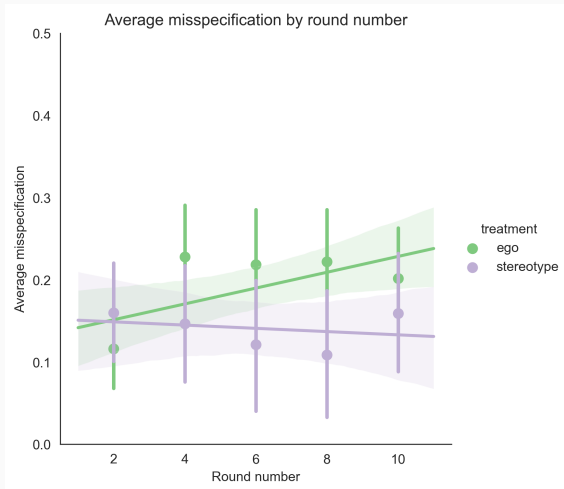


Certainties

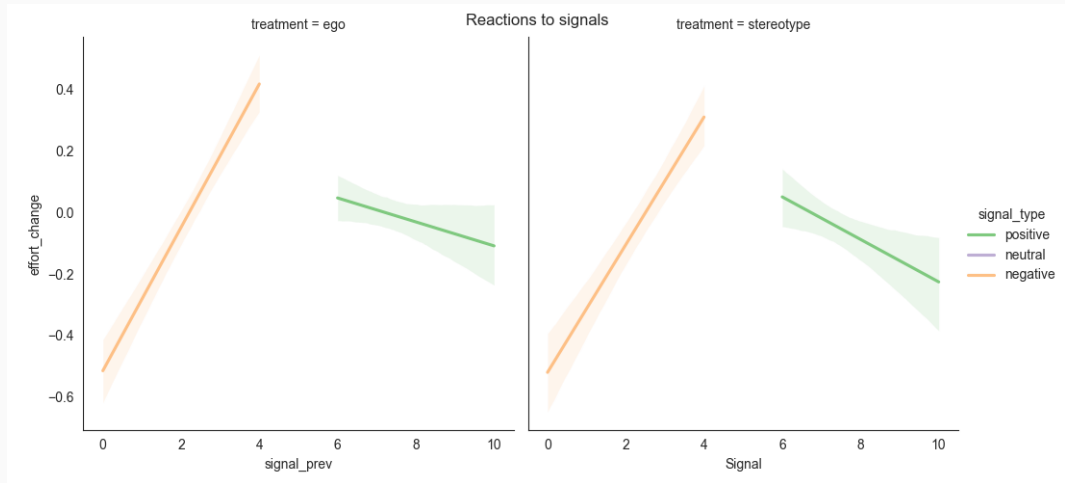


Back


Misspecification changes by treatment



Positive Signals v. Negative Signals



Dogmatic v. Switcher



Bayesian v. Self-Attribution

